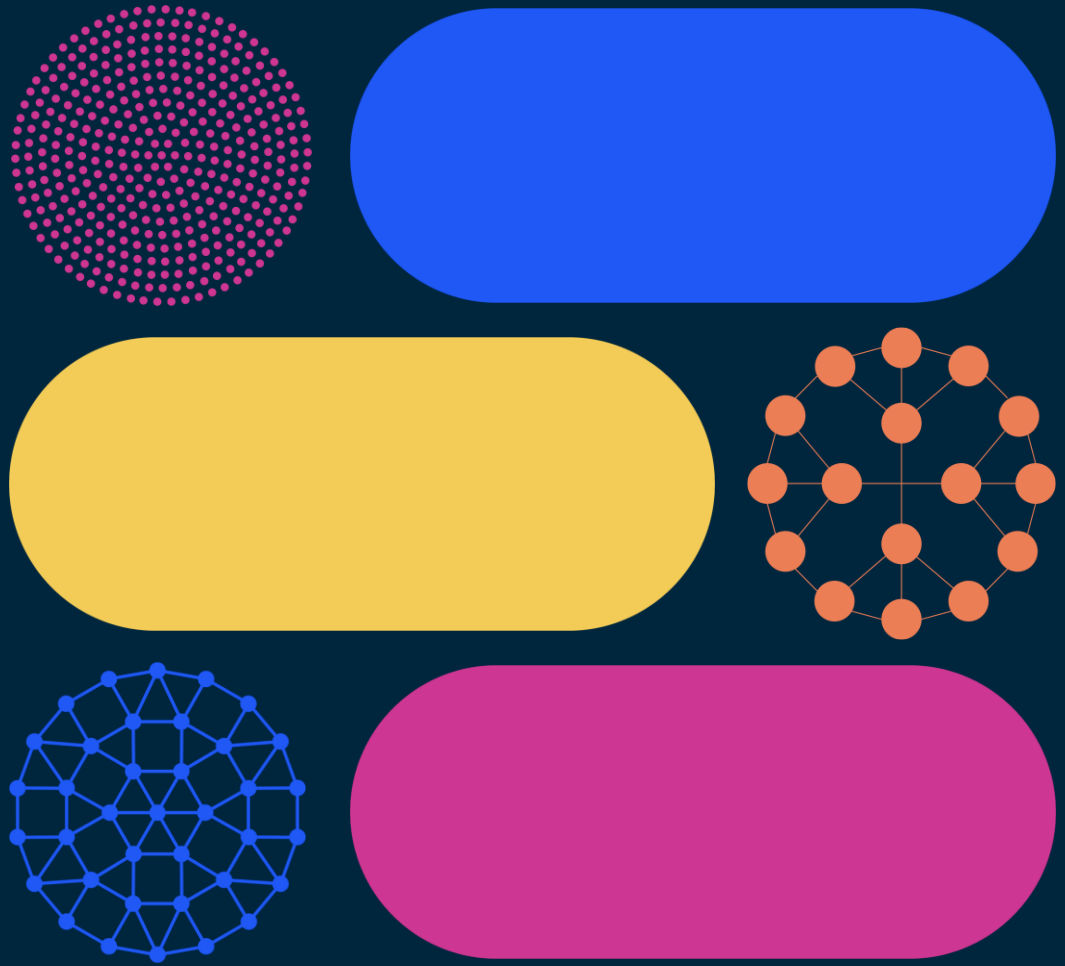

Persistent Environments with Run:ai (using conda/mamba & Jupyter)

Background

A (specific) frequently asked question

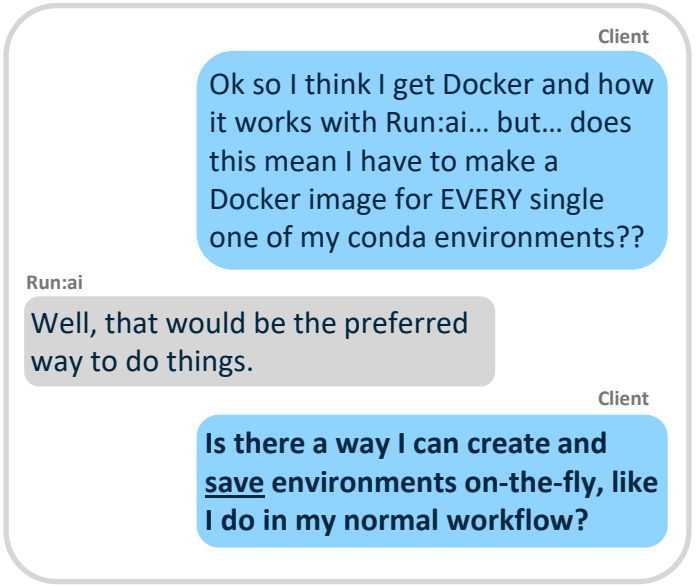

Answer: Yes!

Solution

How do we do this?

Conda background

1. When conda is installed, it creates a directory to keep its files.
2. Whenever a conda environment is created, it is saved inside a directory called “envs,” within the conda directory.

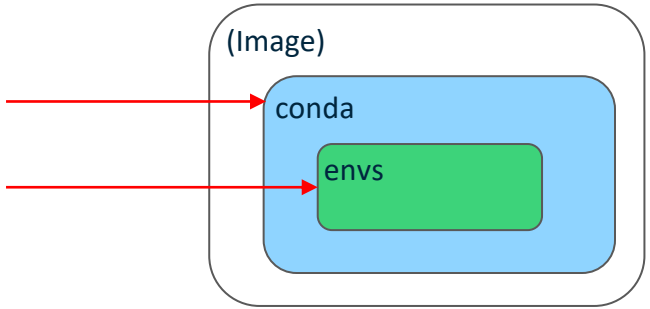

The idea

1. **Create** a **folder** called “**envs**” **in** your **NFS**
2. **Mount** the **NFS** “envs” folder, **to** the location of the **conda environments folder** in your image.

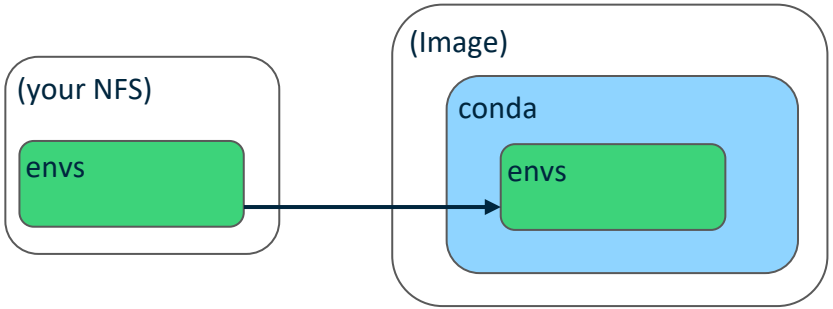

A note on Jupyter

Many users like to have **'tiles' for their environment** on Jupyter

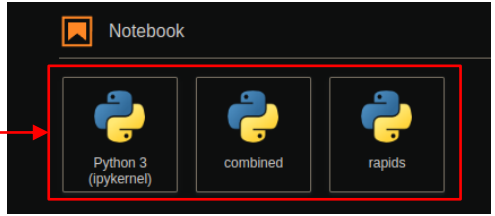

These setting are saved in a folder called 'kernels'

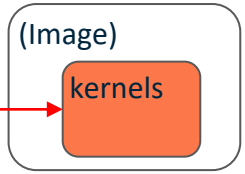

We can **persist** these 'tiles' by:

1. **Creating** a **folder** called "**kernels**" **in** your **NFS**
2. **Mounting** the **NFS** "kernels" folder, **to** the location of the **kernels folder** in your image.

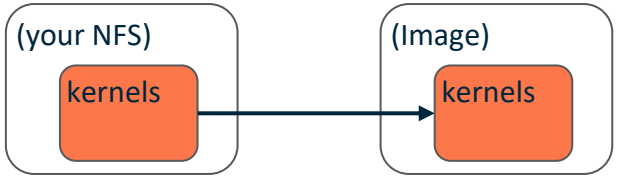

Adding a tile is done with the following sequence of commands, then **refreshing the web browser**

```
mamba create -n {your-env-name} ipykernel {your-env-libs}
conda activate {your-env-name}
python -m ipykernel install --user --name={your-env-name}
```

(If you want to remove the 'tile', this command is used)

```
jupyter kernelspec uninstall {your-env-name}
```

Docker image

The docker image that we will be using throughout this notebook is:
jonathancosme/base-notebook-root*

It is a slightly modified version of the
jupyter/base-notebook
image

You can find more information about jonathancosme/base-notebook-root here:
https://github.com/jonathancosme/jupyter-base-notebook-root

*Please note that all our Jupyter images are based off the official Jupyter Docker images. The official images do not have root privileges by default. In order to successfully mount to the 'kernels' and 'envs' folders in the Jupyter Lab, root privileges are required. We created a custom image that runs root by default. You can see how our image differs from the official image on the github.

Steps

Steps

1. Create ‘envs’ and ‘kernels’ folder in our NFS
2. Mount NFS ‘envs’ and ‘kernels’ folders on an interactive Jupyter lab job
3. Create your environment
4. Test the environment

Step 1

Create ‘envs’ and ‘kernels’ folder in our NFS

1.1 Start a Jupyter Lab interactive session, mounting your NFS folder to /home/jovyan/work

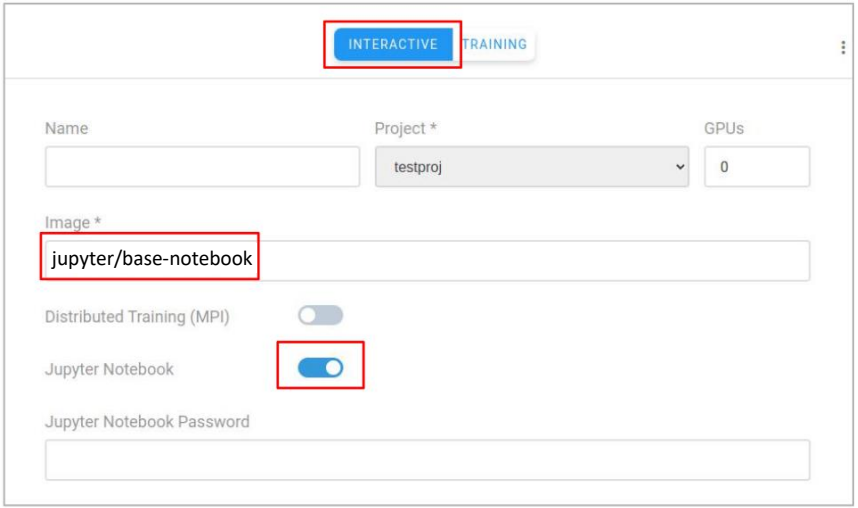

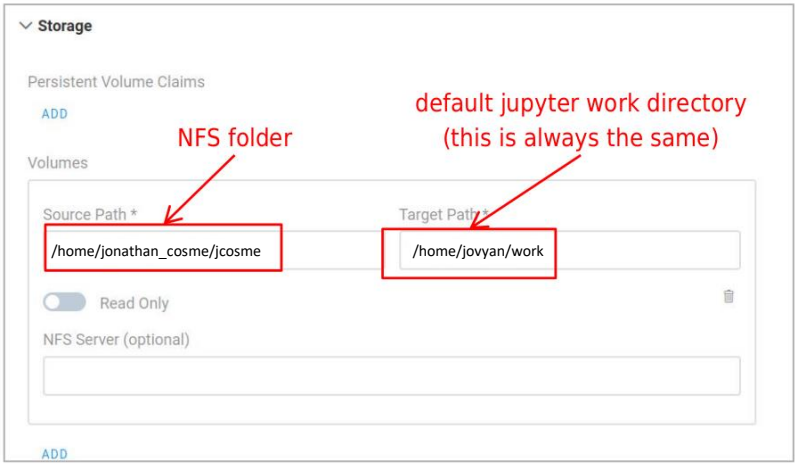

Step 1

Create ‘envs’ and ‘kernels’ folder in our NFS

1.2 Navigate to the ‘work’ directory. You should see your files/folders inside.

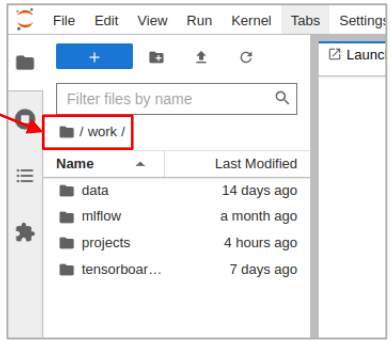

1.3 Select the new folder icon, twice (to create two new folders)

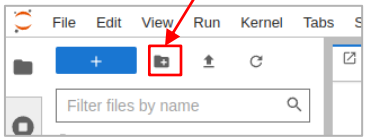

1.4 Rename the folders ‘envs’ and ‘kernels’

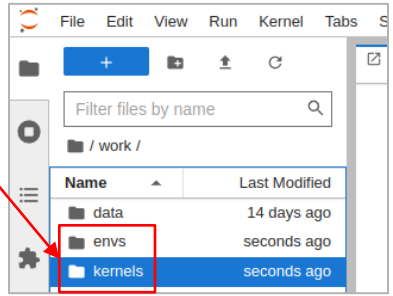

1.5 Delete the job

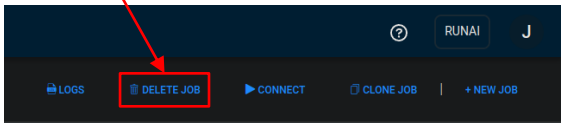

Step 2

Mount NFS ‘envs’ and ‘kernels’ folders on an interactive Jupyter lab job

2.1 Start a Jupyter Lab interactive session using the image **jonathancosme/base-notebook-root** (do NOT submit job yet)

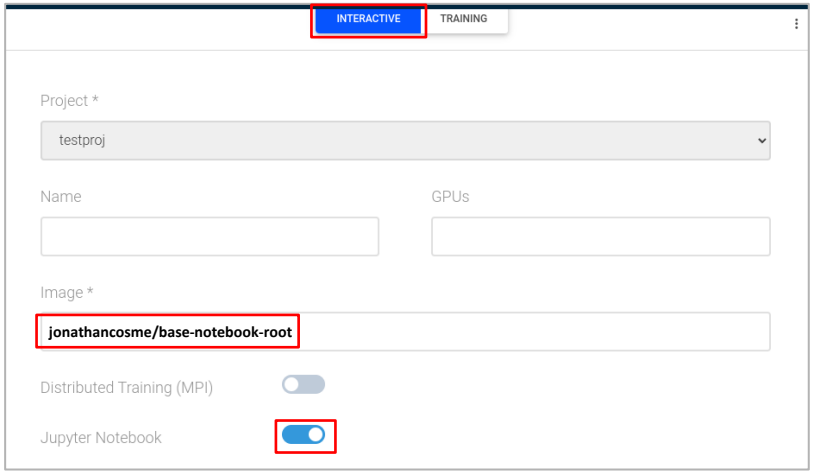

Step 2

Mount NFS ‘envs’ and ‘kernels’ folders on an interactive Jupyter lab job

2.2 mounts NFS to /home/jovyan/work

2.3 mounts NFS ‘kernels’ folder to /home/jovyan/.local/share/jupyter/kernels

2.4 mounts NFS ‘envs’ folder to /opt/conda/envs

2.5 Submit job

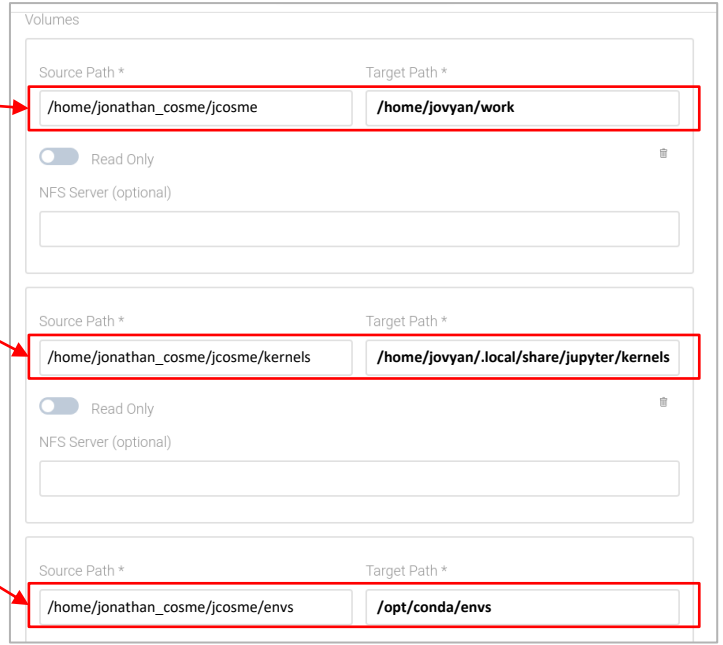

github.com/jonathancosme/runai_persist_envs

Step 3

Create your environment

3.1 start a terminal, after connecting to the Jupyter lab interactive job

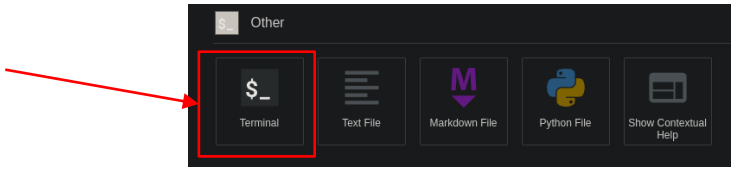

3.2 create an environment using the following format
(ipykernel must always be installed in an environment to create a tile)

```
mamba create -n {your-env-name} -c conda-forge ipykernel {your-env-libraries} -y
```

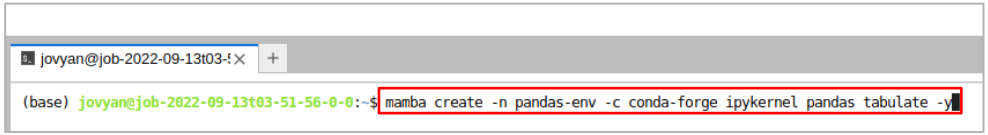

3.3 wait for environment to finish creating.
Warning: depending on your NFS speed, the 'Executing transaction' step can take a long time (~3 minutes for me)

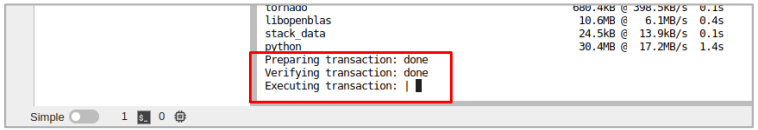

Step 3

Create your environment

3.4 activate the newly created environment using the following format

conda activate {your-env-name}

```
(base) jovyan@job-2022-09-13t03-51-56-0-0:~$ conda activate pandas-env
```

3.5 Create the tile by using the following format

Python -m ipykernel install --user --name={your-env-name}

```
(base) jovyan@job-2022-09-13t03-51-56-0-0:~$ conda activate pandas-env
(pandas-env) jovyan@job-2022-09-13t03-51-56-0-0:~$ python -m ipykernel install --user --name=pandas-env
```

3.6 exit the terminal with the following command

exit

```
(base) jovyan@job-2022-09-13t03-51-56-0-0:~$ conda activate pandas-env
(pandas-env) jovyan@job-2022-09-13t03-51-56-0-0:~$ python -m ipykernel install --user --name=pandas-env
Installed kernelspec pandas-env in /home/jovyan/.local/share/jupyter/kernels/pandas-env
(pandas-env) jovyan@job-2022-09-13t03-51-56-0-0:~$ exit
```

Step 3

Create your environment

3.7 Refresh the Jupyter Lab tab on your browser

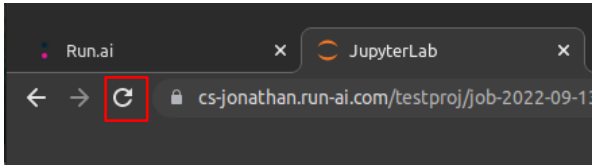

3.8 A tile for your new environment should now be visible

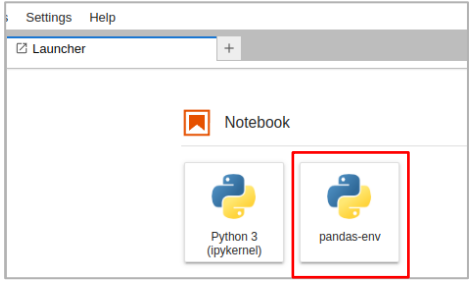

3.9 You can also select the new env is notebooks

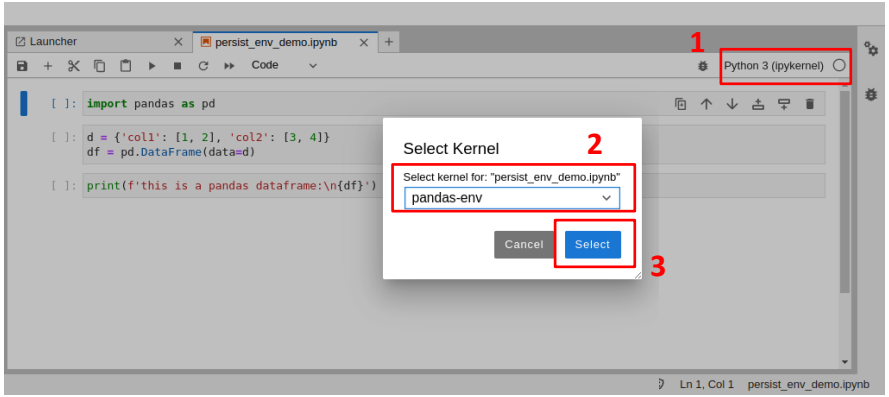

Step 4

Test the environment

Let’s assume I have this python script….

persist_env_demo.py

```python
import pandas as pd
from tabulate import tabulate

d = {'col1': [1, 2], 'col2': [3, 4]}
df = pd.DataFrame(data=d)

tab_df = tabulate(df, headers='keys', tablefmt='psql')

print(f"this is a pandas dataframe:\n{tab_df}")
```

…located here (after mounting my NFS to the Jupyter Lab work directory)

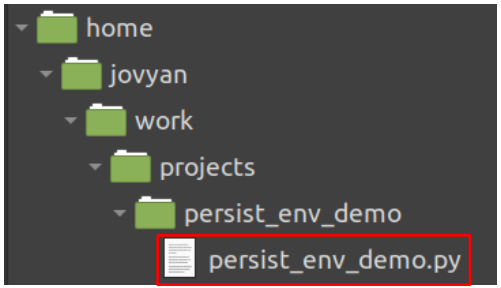

Step 4

Test the environment

This is the runai CLI command I want to run

```
runai submit \
    --project testproj \
    --gpu 0 \
    --job-name-prefix persist-env \
    --image jonathancosme/base-notebook-root \
    --volume /home/jonathan_cosme/jcosme:/home/jovyan/work \
    --volume /home/jonathan_cosme/jcosme/envs:/opt/conda/envs \
    --volume /home/jonathan_cosme/jcosme/kernels:/home/jovyan/.local/share/jupyter/kernels \
    -- conda run -n pandas-env python work/projects/persist_env_demo/persist_env_demo.py
```

Step 4

Test the environment

This is the runai CLI command I want to run

```
runai submit \
    --project testproj \
    --gpu 0 \
    --job-name-prefix persist-env \
1   --image jonathancosme/base-notebook-root \
2   --volume /home/jonathan_cosme/jcosme:/home/jovyan/work \
3   --volume /home/jonathan_cosme/jcosme/envs:/opt/conda/envs \
4   --volume /home/jonathan_cosme/jcosme/kernels:/home/jovyan/.local/share/jupyter/kernels \
5   -- conda run -n pandas-env python work/projects/persist_env_demo/persist_env_demo.py
                                  6
```

1. Use the image jonathancosme/base-notebook-root
2. Mount NFS to Jupyter work directory
3. Mount NFS ‘envs’ folder to Jupyter conda envs folder
4. Mount NFS ‘kernels’ folder to Jupyter kernels folder
5. Run our command (to follow) in the conda environment named ‘pandas-env’
6. Run our python script

Step 4

Test the environment

4.1 Run the runai CLI command. You should see a new job in the job list

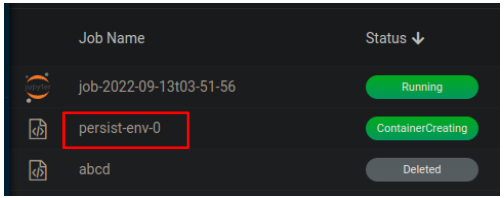

4.2 Wait for the job status to switch to ‘Succeeded’

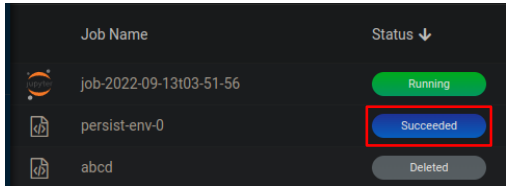

4.3 View the output of the jobs in the Logs tab

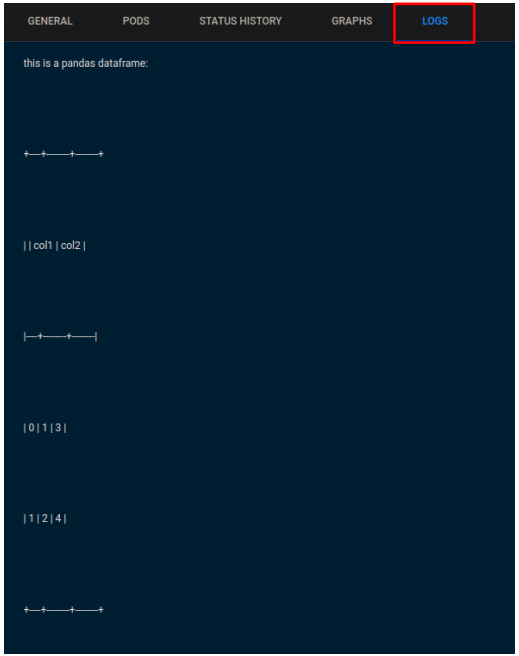

github.com/jonathancosme/runai_persist_envs

Thank you!

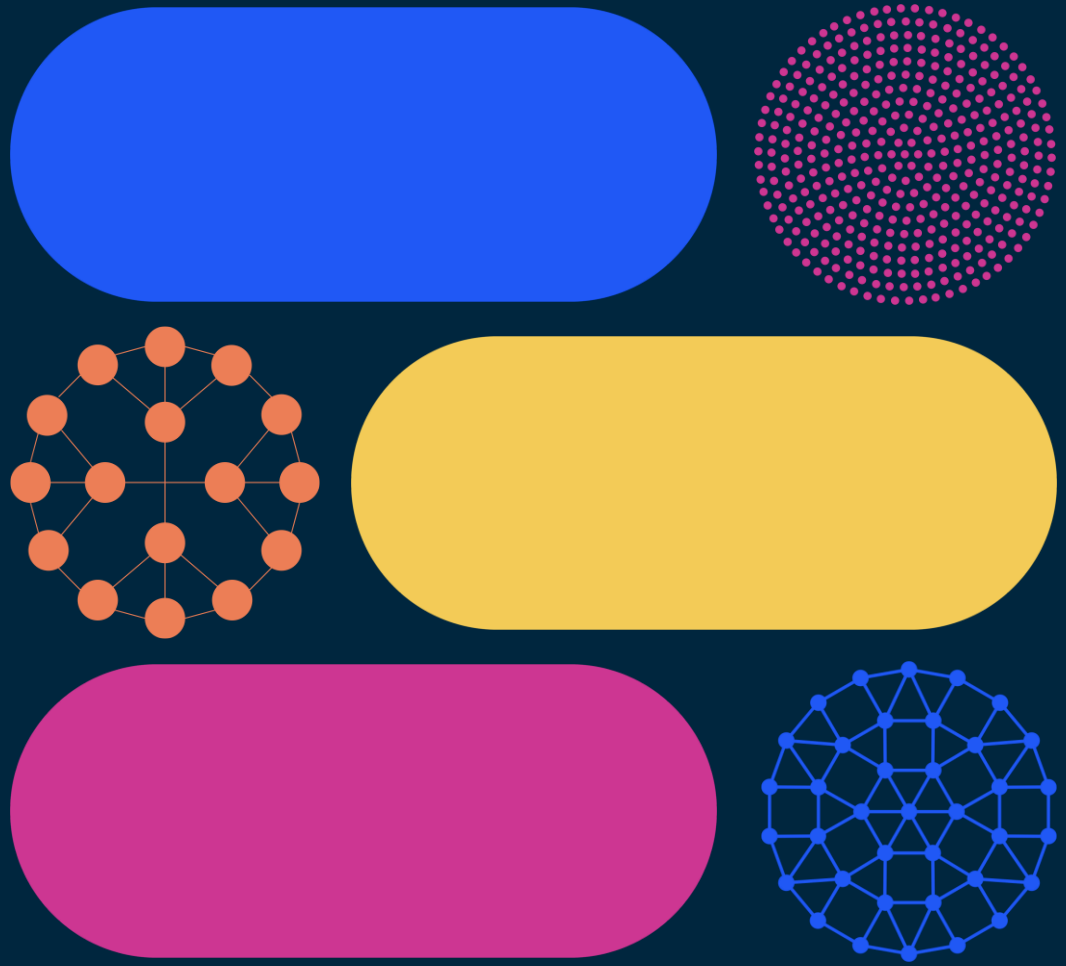

info@run.ai | www.run.ai